Library Research Report (Part 1) Summary
Topic: Responsible AI in the use of generative AI tools
Student: Colton Zachary UCO Username: czachary
Course: CMSC 2123 Date: 02/23/2026

Toward Responsible AI in the Era of Generative AI: A Reference Architecture for Designing Foundation Library: IEEE Xplore
Publication info: IEEE Software, Vol. 41, No. 6, Nov/Dec 2024, pp. 91-100.
Proposes a reference architecture and design guidance for building foundation model based systems responsibly, emphasizing governance and engineering practices to manage risks and impacts of generative AI.

AI Code Generators for Security: Friend or Foe?
Library: IEEE Xplore
Publication info: IEEE Security & Privacy (accepted for a future issue; content final), 2024.
Evaluates AI code generators through a security lens, discussing how they can help or harm security work and what safeguards and evaluation practices are needed for responsible use.

Grow Your Artificial Intelligence Competence
Library: IEEE Xplore

Publication info: Computer, Oct 2024, pp. 145-151 (7 pages).
Argues that responsible adoption of AI (including generative AI copilots) requires building competence in risk management, ethics, security, and vigilant human oversight, with concrete recommendations for individuals and organizations.

Grow Your Artificial Intelligence Competence

Publisher: IEEE Cite This PDF

Christof Ebert ; Ulrich Hemel All Authors

4 Cites in Papers | 613 Full Text Views

Talking about Large Language Models
Library: ACM Digital Library
Publication info: Communications of the ACM, Vol. 67, No. 2, Feb 2024, pp. 68-79.
Explains how large language models work and why their behavior can be hard to predict, highlighting the reliability and safety considerations that motivate responsible deployment and oversight.

Talking about Large Language Models

Author: Murray Shanahan Authors Info & Claims

Communications of the ACM, Volume 67, Issue 2 • Pages 68 - 79 • https://doi.org/10.1145/3624724

Published: 25 January 2024 Publication History Check for updates

210 24,414 PDF/eReader

The Price of Intelligence
Library: ACM Digital Library
Publication info: Communications of the ACM, 2025, pp. 46-53 (8 pages).
Discusses the tradeoffs and risks that come with powerful generative models (e.g., failure modes, misuse, and operational costs), arguing for risk aware and accountable deployment practice.

The Price of Intelligence

Three risks inherent in LLMs.

Authors: Mark Russinovich, Ahmed Salem, Santiago Zanella-Béguelin, Yonatan Zunger | Authors Info & Claims

Communications of the ACM, Volume 68, Issue 9 • Pages 46 - 53 • https://doi.org/10.1145/3749447

Published: 19 August 2025 Publication History

Check for updates

2 14,570

PDF/eReader

To Post or Not to Post: AI Ethics in the Age of Big Tech
Library: ACM Digital Library
Publication info: Communications of the ACM, Vol. 68, No. 8, Aug 2025, pp. 38-49.
Uses current social media and platform dynamics to examine ethical responsibilities and governance questions for AI systems, including generative AI, focusing on accountability, harms, and transparency.

To Post or Not to Post: AI Ethics in the Age of Big Tech

The case of X provides a good example through which to explore the relative merits of descriptive, normative, and action AI ethics.

Author: Henrik Skaug Sætra | Authors Info & Claims

Communications of the ACM, Volume 68, Issue 8 • Pages 38 - 49 • https://doi.org/10.1145/3708339

Published: 15 July 2025 Publication History

Check for updates

0 5,795

PDF/eReader